



Mastering the Game *No Thanks!* with Deep Neural Networks and Tree Search

AMOD 5310H ARTIFICIAL INTELLIGENCE TERM PROJECT

Smars Hu, Trung Kien Ngo, Lya Wang

March 26, 2025



Table of Contents

1 Introduction

► Introduction

► Methodology

► Progress



Motivation

1 Introduction

- *No Thanks!*: popular, multiplayer, turn-based, stochastic card game, centered around risk and resource management.
- *Goal*: Build an AI gaming bot for *No Thanks!*
- *Algorithm*:
 - Inspired by *AlphaGo*, the well-known breakthrough in AI that mastering Go game.
 - Model-based RL: Combining Monte Carlo Tree Search (MCTS) and deep neural network.
 - Upgraded version: *AlphaZero*, no prior knowledge RL.



The Game *No Thanks!*

1 Introduction

- 3-7 players, 33 cards (3 to 35 points), 11 coins (each is worth -1) for each player.
- To win: get *lowest* score.
- Each turn: take the card revealed or spend a coin to pass it; if take the card, you can also take the coins accumulated on it.
- 9 cards are removed from the deck randomly.
- Consecutive cards count only for the lowest-numbered card in the sequence.



AlphaZero Algorithm

1 Introduction

- Generalized reinforcement learning algorithm extending AlphaGo Zero and AlphaGo.
- Trained solely through self-play reinforcement learning.
- A single deep network to estimate both the policy and value functions simultaneously.
- MCTS guided by the policy-value network to perform self-play (PUCT).
- The network is then iteratively updated based on self-play data, with each iteration refining move selection and evaluation.



Contribution

1 Introduction

- To the best of our knowledge, there has been no such algorithm applied to this specific game.
- Explore and extend the applicability of AlphaZero-like algorithms to multiplayer, stochastic games.



Table of Contents

2 Methodology

► Introduction


► Methodology

► Progress



- Game State:** (M, b)

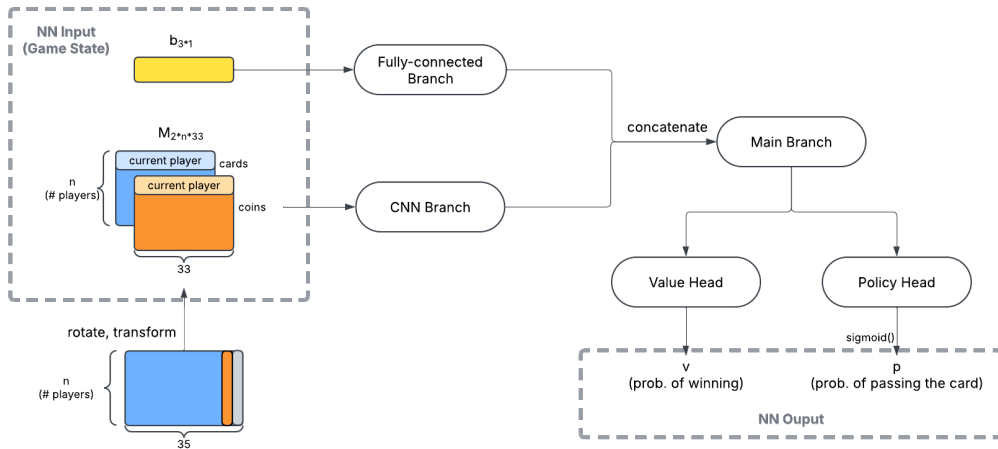
- M : Matrix representing cards and coins holding by each player;
- b : vector: (card in play, coins in play, remaining cards in deck).





Deep Network Architecture

2 Methodology





Monte Carlo Tree Search (MCTS) Self-play

2 Methodology

- Keep track of *Edge*: $(s, a) \rightarrow [N(s, a), W(s, a), Q(s, a), P(s, a)]$.
- *Step 1. Selection*: At each state,

$$a_t^* = \arg \max_a \{Q(s_t, a) + U(s_t, a)\}; \quad (1)$$

where

$$U(s, a) = c_{puct} P(s, a) \frac{\sqrt{\sum_{\alpha} N(s, \alpha)}}{1 + N(s, a)}; \quad (2)$$

- *Step 2. Expansion*: Expand from the root to a leaf (terminal state), where the terminal rewards can be found; we define the terminal rewards to be 1 for winning and 0 for losing.



MCTS Self-play (Ct.d)

2 Methodology

- *Step 3. Back-propagation:* Update $\{N(s, a), W(s, a), Q(s, a), P(s, a)\}$ for each edge in the playout such that

$$N(s, a) = N(s, a) + 1; \quad (3)$$

$$W(s, a) = W(s, a) + z; \quad (4)$$

$$Q(s, a) = W(s, a)/N(s, a). \quad (5)$$

- *Step 4. Choose Action at Root:* After searching, at the root, choose the edge with highest visits; i.e.,

$$a_0 = \arg \max_a N(s_0, a). \quad (6)$$



MCTS Self-play (Ct.d)

2 Methodology

- *Step 5. Update Network and Prior:* Gather self-play data (s, a, q) , where $a = 1$ for PASS, $a = 0$ for TAKE, and $q = Q(s, a)$; Update NN weights θ :

$$\begin{aligned}\theta^* &= \arg \min_{\theta} l(\theta) \\ &= \arg \min_{\theta} \left\{ \|q - v_{\theta}\|^2 - [a \log \mu_{\theta} + (1 - a) \log(1 - \mu_{\theta})] + c_{L2} \|\theta\|^2 \right\}.\end{aligned}\tag{7}$$

And finally update prior:

$$P(s, a) = \mu_{\theta^*},\tag{8}$$

where

$$v_{\theta}, \mu_{\theta} = \text{NN}_{\theta}(s).$$



Performance Evaluation

2 Methodology

Evaluate the performance of our gaming bot by testing it against a mixed group of pure online MCTS agents, rule-based agents, and random agents. The online MCTS agent follows the classic UCT algorithm.



Table of Contents

3 Progress

► Introduction

► Methodology

► Progress



Current Progress

3 Progress

- The code ran properly locally and has been tested for small data (10 rounds training, each round 4 games played);
- The bot is not smart because of insufficient training:
 - NN cannot capture the stochastic nature if it's not in data.
- Challenge: Self-play is very slow.
 - Paralleled;
 - Cut the game after some moves - states closing to the end are not as helpful.



To-Dos

3 Progress

- Bring the code to *Kaggle* and train the model with LARGE amount of self-play data;
- Tune hyperparameters to make the bot smarter;
- More experiments to evaluate the performance;
- Optimize the code.



Q&A

Thank you for listening!
Your feedback will be highly appreciated!