



Mastering the Game *No Thanks!* with Deep Neural Networks and Tree Search

AMOD 5310H ARTIFICIAL INTELLIGENCE TERM PROJECT

Smars Hu, Trung Kien Ngo, Lya Wang

March 22, 2025



Table of Contents

1 Introduction

► Introduction

► Methodology

► Progress



Motivation

1 Introduction

- *No Thanks!*: popular, multiplayer, turn-based, stochastic card game, centered around risk and resource management.
- *Goal*: Build an AI gaming bot for *No Thanks!*
- *Algorithm*:
 - Inspired by *AlphaGo*, the well-known breakthrough in AI that mastering Go game.
 - Model-based RL: Combining Monte Carlo Tree Search (MCTS) and deep neural network.
 - Upgraded version: *AlphaZero*, no prior knowledge RL.



The Game *No Thanks!*

1 Introduction

- 3-7 players, 33 cards (3 to 35 points), 11 coins (each is worth -1) for each player.
- To win: get *lowest* score.
- Each turn: take the card revealed or spend a coin; if take the card, you can also take the coins accumulated on it.
- 9 cards are removed from the deck randomly.
- Consecutive cards count only for the lowest-numbered card in the sequence.



AlphaZero Algorithm

1 Introduction

- Generalized reinforcement learning algorithm extending AlphaGo Zero and AlphaGo.
- Trained solely through self-play reinforcement learning.
- A single deep network to estimate both the policy and value functions simultaneously.
- MCTS guided by the policy-value network to perform self-play (PUCT).
- The network is then iteratively updated based on self-play data, with each iteration refining move selection and evaluation.



Contribution

1 Introduction

- To the best of our knowledge, there has been no such algorithm applied to this specific game.
- Explore and extend the applicability of AlphaZero-like algorithms to multiplayer, stochastic card games.



Table of Contents

2 Methodology

► Introduction

► Methodology

► Progress



Modeling the Game

2 Methodology

- The game is Markovian if all players play with the optimal strategy.

Game State: (M, b)

- M : Matrix representing cards and coins holding by each player;
- b : vector: (card in play, coins in play, remaining cards in deck).

	3, 4, 5, 6, 7, 8, ..., 32, 33, 34, 35, coin, current	
1P	0, 1, 0, ...	3 0
2P		7 0
.		.
.
.		0
kP	1, 0, 0, ...	15 1
.		0
.
nP		4 0



	3, 4, 5, 6, 7, 8, ..., 32, 33, 34, 35, coin, current	
kP	1, 0, 0, ...	15 1
(k+1)P		.
.		.
.
nP		4 0
1P	0, 1, 0, ...	3 0
2P		7 0
.		.
.
(k-1)P		.

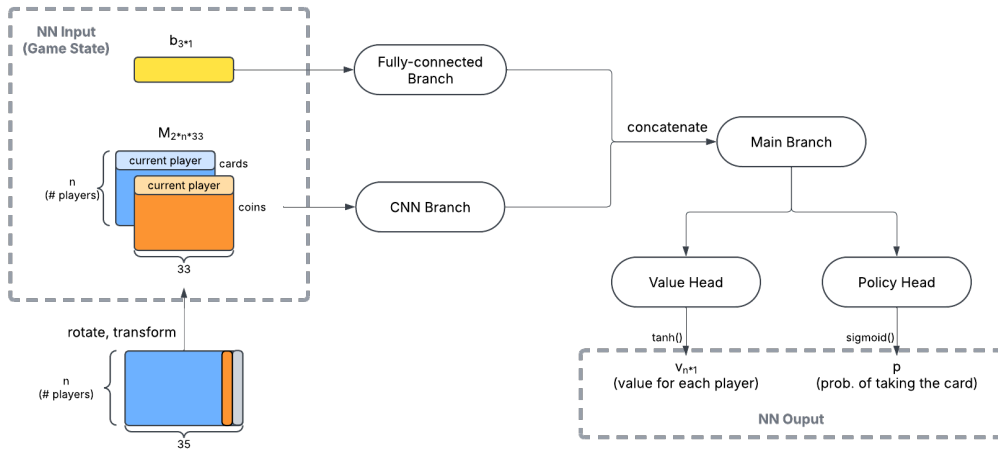


	3, 4, 5, 6, 7, 8, ..., 32, 33, 34, 35	
kP	1, 0, 0, ...	cards
(k+1)P		.
.		.
.
nP	0, 1, 0,
1P		15 15
2P		15
.		.
(k-1)P		.
		.
	4 4	4
	3 3	3
	7 7	7
coins	.	.
	.	.
	.	.
	.	.
	.	.



Deep Network Architecture

2 Methodology





Monte Carlo Tree Search (MCTS) Self-play

2 Methodology

- Keep track of *Edge*: $(s, a) \rightarrow [N(s, a), W(s, a), Q(s, a), P(s, a)]$.
- *Step 1. Selection*: At each state,

$$a_t^* = \arg \max_a \{Q(s_t, a) + U(s_t, a)\}; \quad (1)$$

where

$$U(s, a) = c_{puct} P(s, a) \frac{\sqrt{\sum_{\alpha} N(s, \alpha)}}{1 + N(s, a)}; \quad (2)$$

- *Step 2. Expansion*: Expand from the root to a leaf (terminal state), where the terminal rewards can be found; e.g., \vec{z} such that $z_k = 1, 0.5, 0, -0.5, -1$ for Player k being rank one to five.



MCTS Self-play (Ct.d)

2 Methodology

- *Step 3. Back-propagation:* Update $\{N(s, a), W(s, a), Q(s, a), P(s, a)\}$ for each edge in the playout such that

$$N(s, a) = N(s, a) + 1; \quad (3)$$

$$W(s, a) = W(s, a) + \vec{z}; \quad (4)$$

$$Q(s, a) = W(s, a)/N(s, a). \quad (5)$$

- *Step 4. Update network and prior:* Gather self-play data (s_t, π_t, z_t) , where π_t represent the posterior distribution of actions a_t in edges that stems from s_t , and $\vec{z}_t = \pi_t' Q(s_t, a_t)$; Update NN weights θ :

$$\theta^* = \arg \min_{\theta} l(\theta) = \arg \min_{\theta} \left\{ \|\vec{z} - \vec{v}\|^2 - \pi' \log \mu + c \|\theta\|^2 \right\}. \quad (6)$$

And finally update prior:

$$P(s, a) = (\text{the policy output of}) \text{NN}_{\theta}(s) \quad (7)$$



Performance Evaluation

2 Methodology

Evaluate the performance of our gaming bot by testing it against a mixed group of pure online MCTS agents, rule-based agents, and random agents. The online MCTS agent follows the classic UCT algorithm.



Table of Contents

3 Progress

► Introduction

► Methodology

► Progress



Current Progress

3 Progress

- tba



To-Dos

3 Progress

- Tune hyperparameters to make the bot smarter;
- More experiments to evaluate the performance;
- Optimize the code.



Q&A

Thank you for listening!
Your feedback will be highly appreciated!